

Sean Winslow

(917) 886-1251 | sean.winslow28@gmail.com | [LinkedIn](#) | [Github](#) | seanwinslow.com | Boston, MA

SUMMARY

AI Product Manager who ships agent systems built to keep working. Every release carries eval gates, spend caps, refusal rules, and human review before anything runs. At The Block, launched the company's first sponsored-course B2B revenue line, shipped 3 production Claude Skills behind human review gates, and co-authored the product audit for the incoming CEO. Spent the prior decade running an independent practice inside New York Life, leading media operations and AI workflow integration across 50+ locations. The proof is open source and installable. A governed company OS, a published MCP server, and an agent fleet that runs unattended under eval gates. The systems are new. The judgement isn't.

SELECTED AI WORK

Groundwork: Open-Source Company OS, Governed by Design

- Harness-agnostic operating system a coding agent generates from a structured interview of your company. The generator refuses to invent accountability. Owners, forbidden actions, and retirement conditions come only from human answers, and every generated skill carries a human-written pause and kill condition. Two active competitors shipped adjacent shapes first. Neither ships governance, which is the lane this product owns.
- A zero-dependency validator gates the repo. Changing a rule or a high-risk skill without a matching human-approved proposal is a hard error, while bounded changes auto-apply into an append-only changelog. V1 complete, proven end to end against an adversarially staffed simulated company with a blind-graded scorecard.

Intent-Engineering MCP Server

- 3-tool TypeScript MCP server that audits a natural-language spec against a 9-section template before it reaches a coding agent. Returns a structured rubric on acceptance-criteria testability, scope boundaries, and failure-mode coverage. Scores 23/25 auditing its own spec, and the anti-patterns it caught dogfooding my own 118-skill library are published in the README.
- Live on npm ([@swins/intent-engineering-mcp](#)) and the official MCP registry under a DNS-verified namespace. Failure modes are named and handled. Registry schema drift gets a pinned SDK and an install health check, and a file-disclosure bug I caught myself now routes every disk read through one guarded path. Written scope control shipped it 13 days early.

Code-Brain: Open-Source Agentic Engineering Toolkit

- Open-source Claude Code command center where an agent fleet runs my PM, research, and content workflows unattended on daily and weekly schedules, with structured logging and eval-gated outputs. Agents own decomposition. A human owns judgement.
- Built the measurement layer the honest way. A 10-case eval suite from a 6-mode failure taxonomy mined out of 17 days of production logs, shipped red on purpose at 1/10 and climbed to 7/10. No synthesizer change ships without passing it. A judge layer checks every agent action against an 8-field intent schema on a local model at \$0 per decision, and it fails open because the human gate stays the real control.

Anima: Human + Agent-Fleet 2D Animation Pipeline

- 10-phase pipeline pairing a human director with named agent personas. Agents propose, the human approves, and nothing burns compute until a cost-estimated plan clears a human gate. An LLM-as-judge vision critic scores generated frames at 0.97 precision / 1.00 recall across a 52-case eval suite. First reference implementation shipped Act 1 to a public portfolio.

WORK EXPERIENCE

The Block

Boston, MA (Remote)

Associate Product Manager

November 2025 - May 2026 (role eliminated in a company-wide layoff)

- Stood up The Block's first sponsored-microcourse B2B revenue line (Polymarket × Campus). Owned the PRD across three iterations, shipped the 5-component build, helped close a five-figure inaugural partnership, and templated the motion so the revenue team could repeat it.
- Shipped 3 production Claude Skills (ETF page creation, stakeholder updates, Jira automation), each behind a human review gate, cutting ~6 hrs/week of manual work across the Rev-Ops and P&E teams. Led daily cross-functional P&E standups through a CPO departure and a CEO transition.

- Co-authored the Block Pro 2.0 product audit for the incoming CEO. Benchmarked 9 enterprise data and research platforms, ran 3 stakeholder interviews, and led an 11-risk pre-mortem that surfaced engineering-capacity and renewal-cliff blockers before the roadmap was committed.
- Built RevOps automation from a Salesforce “Closed Won” trigger to auto-created Jira tickets, client intake emails, and Slack routing, removing 7 manual handoffs (~3 hrs per deal).
- Authored the x402 / A2A / MCP agent-economy strategy for Block Pro, translating six monetization patterns into testable product bets.

New York Life Insurance, Multimedia & Design

New York, NY

Media Operations Lead (Independent Contractor)

March 2016 - December 2025

- Retained on a direct 1099 basis for nearly a decade as an embedded operations lead inside a Fortune 100 insurer's in-house creative department.
- Led an 8-person team that rebuilt the enterprise DAM metadata pipeline for 100+ users across 50+ locations and lifted asset discoverability 60%. Introduced LLM-assisted tagging (ChatGPT, Claude, Gemini) as the tooling matured.
- Led the rollout of a SaaS DAM platform across 50+ locations, turning stakeholder requirements into technical documentation in the Atlassian suite.
- Raised media-asset productivity 40% through onboarding and custom training for 100+ users.
- Became the department's go-to AI resource from 2021 onward, the person other teams came to for generative tooling evaluations and recommendations, years before the company had a formal AI function.

EDUCATION

College of Staten Island

Staten Island, NY

B.A. Media Studies

2010-2014

SKILLS

Evals & AI Governance: Golden-set design, LLM-as-judge rubrics, regression suites, failure-mode taxonomies, error analysis, human-in-the-loop deployment, spend governance, agent guardrails and refusal design

Agentic Engineering: Claude Code, Codex, Claude Agent SDK, Claude Skills authoring, MCP servers (TypeScript), agent orchestration, prompt engineering, RAG / local-LLM workflows (Ollama, Qwen-3.5), AI media generation (Nano Banana 2, Veo 3.1, ElevenLabs, Seedance 2.0)

Product: 0-to-1 discovery, PRD authoring, user research, product strategy, go-to-market (GTM), competitive analysis, build-in-public

Tools & Platforms: GitHub, Python, TypeScript, SQLite, Jira, Zapier, Figma, Slack, Obsidian, Anthropic API, NotebookLM

Domains: Agentic infrastructure & developer tools, Enterprise AI & governance, B2B SaaS, Fintech & digital assets, EdTech, AI-assisted media & animation